

Product backlog, sprint backlogs and Definition of Done

Planning counterpart to the delivered portfolio. Every item below is traced to an artefact that exists in this folder. The model file, the 36 forms and the worker sources are the source of truth; where a planning value (owner, estimate, sprint date) is not derivable from an artefact it is marked as a planning value and listed in the Assumptions section at the end.

Read with: `docs/01-gap-audit.md` (the measured starting point), `docs/02-status.md` (what is done and not done), `docs/03-test-record.md` (the execution record), `docs/04-modelling-decisions.md` (why the model is shaped as it is), `docs/05-v2.1-changes.md` (the second release), `docs/06-diagram-engineering.md` (the drawing: how it is routed, and what the third release changed).

1. Product vision and scope statement

Build one runnable BPMN collaboration of the hospital referral, treatment and administration pathway in which every hand-off between the named hospital teams is a real, correlated message; every clinical, administrative and financial decision sits with the correct role; every modelled failure has an owner and an exit; and the whole thing can be regenerated, deployed and demonstrated from the files in one shared workspace.

In scope for the product. The fifteen participants in the case study; the whole pathway from referral receipt to follow-up; the nine hospital processes that carry executable work; the six outside participants as white-box pools of their own since v14 (`docs/11-white-box-external-participants.md`) at the boundary; the nine named exception conditions; compensation for the two resources that are committed and can be released; the clinic-letter escalation ladder; the CAMUNDA Forms and their validation; the external workers that automate 35 job types; and the evidence that all of it runs.

Out of scope for the product (this module). Real integrations with any scheduling, correspondence, treatment, laboratory, imaging or payment provider; production identity, RBAC and persistence; a production database; and any deployment outside a local c8run cluster. See `docs/04-modelling-decisions.md` §7 and `workers/README.md` ("Simplifications and known gaps").

Module learning outcomes used in this document (legend only; the wording of the outcomes is not reproduced in the workspace, see Assumption A7):

Code	Short form used here
LO1	User-centred forms and data capture
LO2	A runnable, correct model and working automation
LO3	Project and test planning; configuration management
LO4	Business behaviour, exceptions and evidence of execution
LO5	Reflective, consistent, well-governed delivery

2. Releases

Two releases exist, matching the two sprints and the two Sprint Reviews.

Release	Name	Files	What it adds	Backlog items
R1	Initial Release (v2.0)	model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn , model/forms/*.form	Measured rebuild: new layout engine, lanes, collapsed subprocesses, exception branches with owners, compensation, capped retry, forms, workers, both run paths	PB-01 – PB-24
R2	Second Release (v2.1)	same filenames, regenerated	Visual compaction only: message-flow thinning, routine end-event merge, lane consolidation; nothing scored removed	PB-25 – PB-29

Definition of the Initial Release (R1). The Initial Release is accepted when: the v2.0 collaboration deploys to a local c8run 8.10.0-alpha5 cluster as 9 process definitions and 36 forms in one batch; the workers build and subscribe to every job type the model defines; the normal path and the unreadable-pack exception path each run to a stable state with incidents: none ; and the layout measures zero sequence-flow overlaps and zero message-flow overlaps. All five conditions are evidenced in docs/03-test-record.md §C, §D, §E and §H.

3. Priority scheme

MoSCoW is used, and within each band items are ranked 1 = first. Rank drives sprint selection; MoSCoW drives the release decision.

Band	Meaning for this product
Must	Without it the release is not accepted (see §2)
Should	Required for the 70%+ rubric band; carried as a committed stretch item
Could	Improves the portfolio; carried only if the Must and Should items are done
Won't (this release)	Explicitly deferred, with the reason recorded

4. Product backlog

Estimates are in story points (SP) and ideal days (d). Owners are the five-person team: **A** model lead, **B** layout and visualisation lead, **C** workers lead, **D** forms and UX lead, **E** test and configuration lead. Every item has a first owner (accountable) and a second owner (reviewer and cover).

4.1 Sprint 1 items — delivery of the Initial Release (R1)

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-01	1	Must	Measure v1.0 against every 70%+ rubric descriptor and list the gaps	LO3, LO5	5	1.5	E	A	Every descriptor in the gap audit has either a measurement or an artefact reference; every gap is numbered and actionable	—	Done	docs/01-gap-audit.md ; tools/analyse_layout.py
PB-02	2	Must	One BPMN collaboration with the 15 case-study participants; 9 executable hospital processes and 6 outside participants	LO2, LO4	8	2	A	E	15 bpmn:participant elements with the case-study names; 9 processes with isExecutable="true" ; originally the 6 external bodies were black-box pools, superseded in v14 — all 15 now carry a process and every outside participant is connected to another pool	PB-01	Done, extended in v14	model file participants; docs/04-modelling-decisions.md §1; docs/11-white-box-external-participants.md
PB-03	3	Must	Correlated message hand-off for every cross-pool transfer	LO2, LO4	8	2	A	C	68 message throw events and 10 catch events; every throw carries a correlationKey input mapping; message names are unique	PB-02	Done	model file; docs/03-test-record.md §B1; docs/04-modelling-decisions.md §2
PB-04	4	Must	Lanes inside pools that cover more than one desk	LO2	5	1.5	A	B	Every pool that covers more than one sub-role carries named lanes; participant names unchanged	PB-02	Done	14 lanes, e.g. consultants_L0 / consultants_L1 ; docs/03-test-record.md §A7
PB-05	5	Must	Orthogonal layout engine: grid columns/rows, one track per flow, dedicated message channel	LO2, LO4	13	4	B	A	Sequence-flow overlap 0 pairs, message-flow overlap 0 pairs, 0 diagonal px; both hospital-patient-pathway-v2.svg and .png regenerate	PB-02	Done	tools/layout_engine.py ; docs/03-test-record.md §A1-A3; diagram/

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-06	6	Must	Contain repetitive administration in collapsed sub-processes	LO2	8	2	A	B	3 collapsed sub-processes, each with a <code><bpmn:documentation></code> note; their steps remain in the XML and executable	PB-05	Done	SUB_Secretaries_Dispatch , SUB_CallHandling_Contact , SUB_Pathway_Report ; docs/03-test-record.md §A8
PB-07	7	Must	Every exception branch has an owner and an exit	LO4	13	4	A	C	All 9 <code>bpmn:error</code> codes reachable from 18 boundary events; no exception branch terminates without an owner task, a retry, a timer or an escalation	PB-03	Done	docs/04-modelling-decisions.md §4; docs/03-test-record.md §H1
PB-08	8	Must	Clinic-letter delay timers and the three-rung escalation ladder	LO4	8	2	A	E	Boundary timers <code>P7D</code> and <code>R/PT168H</code> present; <code>letter.escalation-admin-manager</code> and <code>letter.escalation-higher-management</code> message chain present	PB-07	Done	model boundary events; docs/03-test-record.md §H1
PB-09	9	Must	A Camunda Form on every user task	LO1	8	2.5	D	A	39/39 user tasks carry <code>zeebe:formDefinition</code> ; 36 distinct forms exist and every bound form is defined	PB-02	Done	model/forms/ (36 .form); docs/03-test-record.md §H1, §H5d
PB-10	10	Must	Role-based candidate group on every user task	LO1, LO5	5	1.5	D	E	39/39 user tasks carry <code>zeebe:assignmentDefinition</code> with a candidate group matching the pool's role	PB-09	Done	docs/03-test-record.md §H1, §H5d
PB-11	11	Must	Validation rules on the forms	LO1	5	1.5	D	B	Required, length and numeric-range rules present on the fields that need them; long text is bounded	PB-09	Done	172 <code>validate</code> blocks across the 36 forms; docs/01-gap-audit.md §C2
PB-12	12	Should	Accessibility: field descriptions on every field, and grouping on long forms	LO1, LO5	5	1.5	D	E	Every field has a description; forms longer than ~8 fields are split by group heading	PB-09	Partly done — 71 descriptions across 230 fields; 0 <code>group</code> components emitted	form JSON census; docs/01-gap-audit.md §C3

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-13	13	Must	One external worker per job type the model uses	LO2, LO4, LO5	21	6	C	A	Every job type in the model has a subscribing worker; build green; subscriptions observed on the live gateway	PB-02	Done	workers/src/.../tasks/ (36 classes); workers/README.md ; docs/03-test-record.md §E1-E2
PB-14	14	Must	Simulated external services with documented forced-failure switches	LO4	8	2.5	C	E	The 5 simulated job types can each be forced to succeed or fail without editing code, through application.yaml or an environment variable	PB-13	Done	SimulatedFailureRegistry ; workers/src/main/resources/application.yaml ; workers/README.md
PB-15	15	Must	Duplicate suppression for bookings, charges and hand-offs	LO2, LO4	8	2	C	B	Repeating a booking returns the same reference; a duplicate charge is detected before payment; the correlation-key purpose tag prevents a duplicate hand-off	PB-03, PB-13	Done	HospitalStore ; workers/README.md ; docs/04-modelling-decisions.md §2
PB-16	16	Must	No storage or emission of complete card data	LO5	3	1	C	E	No payment worker reads or writes a card-number, CVV or expiry variable; only a payment reference and status are returned	PB-13	Done	workers/README.md job types 13 and 24 ("Never emits card data")
PB-17	17	Must	Audit-trail workers that do not overwrite clinical variables	LO4, LO5	5	1.5	C	A	Clinical and financial decisions produce an audit reference and timestamp; the audit workers refuse a blank patientRef and write no clinical variable back	PB-13	Done	workers/README.md job types 2 and 3
PB-18	18	Must	Compensation for the two committed resources	LO2, LO4	13	4	A	C	A compensation boundary event and isForCompensation handler on the appointment series and on the cycle booking; a compensation throw on the path that gives up; handlers typed treatment.release-series and treatment.release-cycle-booking	PB-07, PB-13	Done for modelling; not triggered live	model TRT_Bnd_CompSeries , TRT_Bnd_CompCycle , TRT_Comp_ReleaseSeries , TRT_Comp_ReleaseCycle ; docs/04-modelling-decisions.md §5

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-19	19	Must	Counted retry with a cap and an escalation for unavailable external capacity	LO4	8	2	A	C	A counted retry capped at 3; on exhaustion the provisional series is released and the case is handed to the pathway team, then retried in three days	PB-07	Done	TRT_Auto_RecordRetry (treatment.record-capacity-retry); docs/04-modelling-decisions.md §4
PB-20	20	Must	Deploy the model and all forms in one batch, repeatably	LO5	5	1.5	E	C	One command deploys 9 process definitions and 36 forms with no error, from a clean checkout, against <code>http://localhost:8080</code>	PB-05, PB-09	Done	tools/deploy.sh ; docs/03-test-record.md §C1, §H2a
PB-21	21	Must	Normal-path run evidence	LO4	5	1.5	E	C	Eleven steps reached; incidents: none reported at the end of each stage; log retained	PB-20	Done	screenshots/run-happy-path.log
PB-22	22	Must	Exception-path run evidence	LO4	5	1.5	E	A	REFERRAL_PACK_UNREADABLE is thrown by the worker, caught by SEC_Bnd_DocumentsUnreadable , and the case joins the missing-information loop with no incident	PB-20	Done	screenshots/run-exception-path.log
PB-23	23	Must	Prove every form renders	LO1, LO5	3	1	D	E	All 36 forms import and render in @bpmn-io/form-js with no unsupported field type	PB-09	Done	screenshots/camunda-forms-rendered.png ; docs/03-test-record.md §F4
PB-24	24	Must	Written test record and modelling decisions document	LO3, LO5	8	2.5	E	A	Expected vs actual for every check run; every defect recorded; limitations stated before a reviewer finds them	PB-21, PB-22	Done	docs/03-test-record.md , docs/04-modelling-decisions.md

4.2 Sprint 2 items and deferred work

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-25	25	Must	Preservation audit as executable checks	LO3, LO5	8	2.5	E	B	Parsed checks over pools, user tasks, forms, candidate groups, job types, gateway types, error codes, timers and Camunda 7 attributes; all pass before a rebuild is accepted	PB-24	Done	tools/verify_preservation.py – 62 checks, 62 passed; docs/03-test-record.md §H1
PB-26	26	Must	Thin the drawn message flows to one line per pool pair, keeping exception hand-offs	LO2	8	2.5	B	A	Message flows 66 → 31; dashed-line length 421,424 → 174,820 px; both run paths still clean afterwards	PB-25	Done	docs/05-v2.1-changes.md §1; docs/03-test-record.md §H3
PB-27	27	Must	Merge routine end events, never exception ends	LO2	5	1.5	B	E	End events 60 → 32; <code>_exception_ends()</code> keeps every exception end distinct	PB-26	Done	docs/05-v2.1-changes.md §2
PB-28	28	Should	Consolidate single-step lanes into a neighbour lane of the same pool	LO2	3	1	B	A	Lanes 20 → 14; no participant renamed or removed	PB-26	Done	docs/05-v2.1-changes.md §3
PB-29	29	Must	Re-validate in the Modeler and export the diagram and per-pool sections	LO4, LO5	8	2.5	B	E	Camunda Modeler 5.51.0 reports 0 errors on Camunda 8.10 (alpha); SVG and PNG regenerate; 15 section crops refreshed	PB-26	Done	screenshots/camunda-modeler-v2.1-status-bar.png ; tools/render_diagram.sh ; tools/export_sections.py ; diagram/sections/
PB-30	30	Should	Drive the three branches that are modelled but not run: declined payment, confirmation lost, three-month escalation	LO4	8	2.5	E	C	Each branch driven to its end state on the live engine, with a retained log and an Operate screenshot	PB-20	Carried over — not run	docs/03-test-record.md §D3, §G, §H6e; docs/02-status.md ("Not done")

ID	Rk	Pri	Item	Why (LO)	SP	d	1st	2nd	Acceptance criteria	Depends on	Status	Evidence
PB-31	31	Should	Trigger compensation in a live run	LO4	5	1.5	C	E	A compensation handler completes inside a real process instance, visible in Operate	PB-18, PB-20	Carried over — not run	docs/03-test-record.md §D3
PB-32	32	Could	Fix the c8run Tasklist authentication so the forms draw in Tasklist	LO1, LO5	3	1.5	E	D	The Tasklist Task tab renders a v2.0 form for a claimed task	PB-20	Carried over — Tasklist gets 401 from /v2/authentication/me ; environmental	docs/03-test-record.md §F5; docs/02-status.md
PB-33	33	Could	Reduce message start events to ≤ 2 per pool by changing the message protocol	LO2	13	4	A	C	Every hand-off publishes <code>Msg_<pool></code> plus a <code>requestType</code> discriminator and each pool has one start event feeding a dispatch gateway	PB-03	Not started — deferred, see docs/05-v2.1-changes.md §4	arithmetic in docs/05-v2.1-changes.md §4
PB-34	34	Could	Crossing-reduction pass	LO2	8	2.5	B	A	Sequence-flow crossings reduced from 197 without reintroducing an overlap	PB-27	Done in R3 — 94 crossings (–52%), 0 clipped lines, 0 overlaps held	docs/06-diagram-engineering.md §3.1, §3.4
PB-35	35	Should	Readable terminals and hand-offs: one end event per terminating path, and dashed lines cut to the hand-offs that carry an outcome	LO2, LO4	5	1.5	B	D	Every end event has exactly one incoming sequence flow; no end event unreachable; every exception hand-off still drawn; no more than 15 message flows drawn	PB-27	Done in R3 — 66 end events, none merged; 15 drawn message flows. v14 raises the message-flow ceiling to 22 to connect the six outside participants	docs/06-diagram-engineering.md §3.5, §3.6; docs/11-white-box-external-participants.md
PB-35	35	Won't (this release)	Replace the five simulated external services with real provider APIs	LO4	—	—	C	E	Not applicable — out of scope for a local portfolio release	PB-14	Won't (this release)	workers/README.md ("Simplifications and known gaps")

5. Sprint backlogs

Sprint lengths, stand-up dates and review dates are planning values; no calendar record exists in the workspace (Assumption A3).

5.1 Sprint 1 — to Sprint Review 1, delivering the Initial Release (R1)

Sprint Goal (agreed at planning).

Replace the v1.0 drawing with a measured layout engine and close the two behavioural gaps the gap audit found — exception branches with an owner and compensation for committed resources — then prove the whole collaboration runs on a live engine with no incidents.

Selection and control. All 24 Sprint 1 items entered the sprint; PB-01 to PB-05 were completed first because every later item depends on the structure and the layout, and the sprint was deliberately not expanded when the forms defect (§B of `docs/03-test-record.md`) was found. PB-12 was descoped during the sprint to descriptions only; the grouping half is retained as unfinished.

Sprint 1 selection	Items	State at Sprint Review 1
Analysis and structure	PB-01, PB-02, PB-03, PB-04, PB-06	Completed
Drawing	PB-05	Completed
Behaviour	PB-07, PB-08, PB-18, PB-19	Completed for modelling; PB-18 not triggered live
Forms	PB-09, PB-10, PB-11, PB-12	PB-09/10/11 completed; PB-12 partly done
Workers	PB-13, PB-14, PB-15, PB-16, PB-17	Completed
Release and evidence	PB-20, PB-21, PB-22, PB-23	Completed
Documentation	PB-24	Completed

Completed at Sprint Review 1 (22 of 24). PB-01 – PB-11, PB-13 – PB-17, PB-18 (modelling only), PB-19 – PB-24. PB-12 was partly done; PB-18 was not triggered live.

Carried over out of Sprint 1.

Item	Carried because	Where it went
PB-12 (grouping half)	The field-description half was delivered; grouping was not implemented and no <code>group</code> component is emitted in any form	Not selected in Sprint 2; recorded as an open gap in <code>docs/08-test-plan.md</code> and <code>docs/10-sprint-reviews-and-feedback.md</code>
PB-18 (live trigger)	Compensation is modelled and the handlers exist; no live run exercised it	Re-framed as PB-31 in Sprint 2, still not run

5.2 Sprint 2 — to Sprint Review 2, delivering the second release (R2)

Sprint Goal (agreed at planning).

Reduce the drawn density of the shipped model without removing a single scored element, and re-prove the behaviour afterwards on the rebuilt file.

Selection and control. The committed set was PB-25 – PB-29, with PB-30 – PB-34 held as stretch items behind the Must items. PB-25 (the executable preservation audit) was built first so that every later change could be proved non-destructive. PB-30 – PB-34 were not started, and no committed item was dropped.

Sprint 2 selection	Items	State at Sprint Review 2
Committed	PB-25, PB-26, PB-27, PB-28, PB-29	Completed
Stretch	PB-30, PB-31, PB-32, PB-33, PB-34	None started
Deferred by decision	PB-35	Won't (this release)

Completed at Sprint Review 2 (5 of 5 committed). PB-25 – PB-29.

Carried over out of Sprint 2. PB-30, PB-31, PB-32, PB-33, PB-34, PB-35. Each carry-over keeps its MoSCoW band and its owner; the reason is recorded in the table in §4.2 and repeated in `docs/10-sprint-reviews-and-feedback.md`.

6. Definition of Done

An item is Done only when every applicable line below is true. The gates are the real commands in `tools/`; a claim without a gate result is not Done.

#	DoD condition	How it is proved	Applies to
DoD-1	The generator accepts the change: <code>python3 tools/build_v2.py</code> exits 0 with no check failure	Build output; the check refuses unbound forms, shared grid cells and mixed boolean/text comparisons	Model, forms
DoD-2	Layout gates hold: sequence-flow and message-flow overlaps are 0 pairs; no diagonals	<code>python3 tools/analyse_layout.py model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn</code>	Drawing
DoD-3	Nothing scored was lost by a regeneration: 67 parsed checks pass	<code>python3 tools/verify_preservation.py</code>	Any model change after R1
DoD-4	The model parses clean in the Modeler's own parser: 0 warnings, 0 unresolved references, 0 dangling flows, 0 user tasks without a form or candidate group, 0 service tasks without a job type	<code>node tools/validate_model.js model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn</code>	Model
DoD-5	The model opens in Camunda Modeler 5.51.0 with 0 errors on Camunda 8.10 (alpha)	Modeler status bar, screenshot retained	Release
DoD-6	The deployment is accepted: 9 process definitions and 36 forms in one batch	<code>bash tools/deploy.sh</code> output	Release

#	DoD condition	How it is proved	Applies to
DoD-7	The workers build and subscribe to every model job type	<code>mvn -q -o clean package</code> ; start-up subscription log	Workers
DoD-8	Both scripted paths reach their expected end state with <code>incidents: none</code>	<code>python3 tools/demo_scenario.py</code> and <code>--exception</code> ; logs retained	Release
DoD-9	Every form imports and renders in <code>@bpmn-io/form-js</code>	Render census, <code>TOTAL=36 FAILED=0</code> ; screenshot retained	Forms
DoD-10	Expected result, actual result and evidence file are written down	The test record is updated in the same change	Every item
DoD-11	The second owner has read the change and the evidence, and their name is against the item	Hand-over log in §7	Every item
DoD-12	Any limitation introduced is stated in <code>docs/04-modelling-decisions.md</code> §7 or <code>docs/02-status.md</code>	Document review at the sprint review	Every item that leaves a gap

7. How the two owners work, and the hand-over log

The pattern. The first owner is accountable: they make the change, run the gates that apply to it, and write the evidence line. The second owner is the reviewer and the cover: they reproduce at least one gate result themselves, read the evidence rather than the claim, and are expected to be able to take the item over within one working session. No item is Done with only one owner's name on it (DoD-11). Where the two owners disagree, the item stays In progress and the disagreement is raised at the stand-up; it is not resolved by the first owner alone. For cross-boundary items the second owner is deliberately chosen from the other side of the boundary — the forms items are second-owned by the test lead, the worker items by the model lead — so that a claim of integration is checked by someone who did not build the integration.

Cadence of hand-over. A hand-over is an explicit act, not an email: the artefact moves into the shared workspace, the gate output is pasted into the test record, and the row in the log below is closed in the same session. The log is the record of who held an item and when.

HO	From → To	Artefact handed over	Trigger	Accepted how	Status
HO-01	A → B	<code>tools/spec_v2.py</code> structure (pools, lanes, nodes, flows)	Structure frozen for Sprint 1	B builds the layout from the spec and returns the overlap measures	Closed
HO-02	B → A	Layout engine geometry and the <code>analyse_layout</code> output	After the first zero-overlap run	A confirms every node and flow is placed and no element moved between lanes	Closed
HO-03	A → C	Model job types and error codes	Before worker work	C confirms every service task type and every BPMN error has a worker behaviour	Closed (three v2.0 types added afterwards, logged as HO-08)

HO	From → To	Artefact handed over	Trigger	Accepted how	Status
HO-04	D → E	36 <code>.form</code> files	Before the render census	E runs the form-js census and reports failures	Closed — first run failed, see <code>docs/03-test-record.md</code> §F4
HO-05	E → D	Form render failure report (<code>date</code> field type)	Same day as HO-04	D fixes <code>tools/forms_spec.py</code> and the build re-emits the forms	Closed
HO-06	E → C	Deployment result (9 definitions, 36 forms)	Before the path runs	C runs the paths and returns the incident count	Closed
HO-07	C → E	Path logs	End of Sprint 1	E signs off the expected-vs-actual table	Closed
HO-08	A → C	Three v2.0 job types (<code>treatment.release-series</code> , <code>treatment.release-cycle-booking</code> , <code>treatment.record-capacity-retry</code>)	After compensation and capped retry were modelled	C confirms the three workers build and register	Closed
HO-09	B → E	v2.1 rebuilt model	Start of the compaction pass	E runs the 62-check preservation audit before anything else is accepted	Closed
HO-10	E → A	Preservation result and Modeler status bar	End of Sprint 2	A confirms the release note and the carried-over list	Closed

8. Agile working notes

Cadence. Two timeboxed sprints. Each sprint runs a graded stand-up at its midpoint and a graded Sprint Review at the end, so the demonstration is rehearsed once before it is graded. Planning is a single short session at the start of each sprint; the sprint backlog is frozen after it except for a removal, which must be agreed by both owners and recorded in §5.

Stand-up format (timeboxed, 10 minutes). Each of the five members answers in turn, with the artefact open rather than from memory: (1) which backlog item did you move, and which gate output proves it; (2) which item will you move next; (3) what is blocking you, and from whom. Blockers are written into the feedback action log in `docs/10-sprint-reviews-and-feedback.md` with a first and second owner before the stand-up ends. Status is only ever one of: Not started, In progress, Blocked, Done (per DoD), Carried over.

Review format (timeboxed). For each demonstrated item: the sprint goal, the increment actually shown, the acceptance evidence file, what was not finished, the decisions taken, and the limitations stated. The review runs the live path, not a slide. Anything not shown at the review is not counted as Done at that review even if it is finished afterwards.

Retrospective format (timeboxed, after each review). Three columns — what went well, what did not, and one improvement action. The action must name a first owner, a second owner and the sprint it lands in; an action with no owner is not recorded. The retrospectives are in `docs/10-sprint-reviews-and-feedback.md` §3 and §6.

Continuous integration of contributions. The workspace is integrated after every accepted item, not at the deadline: the generator is re-run, the layout gate and the preservation gate are executed, and the test record is updated in the same change. The deterministic build is the integration mechanism — `build_v2.py` is documented as producing byte-identical output for the same inputs (`docs/05-v2.1-changes.md` §5), so any unexplained diff is a real change to be reviewed rather than an accident of the tooling.

9. Assumptions

1. **A1 — Team identity.** The five members are referred to as team members A–E with the roles in §4. No real names appear in any artefact in the workspace, or in this document.
2. **A2 — Estimates.** Story points, ideal days and ranks are planning values authored for this document. No estimate exists in any artefact in the workspace.
3. **A3 — Dates.** Sprint length, stand-up dates and review dates are not recorded anywhere in the workspace. They are treated here as planning values for two timeboxed sprints.
4. **A4 — Repository.** There is no `.git` directory or other version-control metadata in the folder (checked). The shared workspace is therefore the folder itself, and the configuration-management plan in `docs/09-risk-contingency-and-config-management.md` describes the scheme as it exists rather than assuming a VCS.
5. **A5 — Job-type counts.** The model defines 47 service tasks across **35 distinct job types**, plus `publish-message` on the 68 message throw events. The worker sources contain 36 task classes producing **37 job-type subscriptions** (35 model types, one alias `pathway.find-overdue-clinic-letters`, and `publish-message`), which reconciles the "37 subscriptions" figure in `docs/03-test-record.md` §E2 and §H2b. The table in `workers/README.md` documents only 32 domain types and its summary line ("33 job workers = 1 + 32") is stale by three classes; this is recorded as a documentation defect in `docs/08-test-plan.md` §9.
6. **A6 — Diagram document (closed in R3).** The earlier `docs/06-visual-compaction.md` described a compaction that was never in the shipped model. In the third release that file was **deleted** and replaced by `docs/06-diagram-engineering.md`, which describes the drawing that is actually shipped and records the abandoned compaction as a rejected experiment with its measurements. The model file remains the source of truth; the drift is closed rather than carried.
7. **A7 — Learning-outcome legend.** The full outcome wording is not reproduced in the workspace; the short forms in §1 are inferred from the descriptor/LO mapping used in `docs/01-gap-audit.md`.
8. **A8 — Release naming.** "Initial Release" is R1 (v2.0) and the "second release" is R2 (v2.1). Both filenames are the same in the folder; the distinction is the content and the deployment version, not the path.