

Risk, contingency and configuration management

Risk register, demonstration contingency and rollback, and the configuration-management plan for the delivered portfolio. Every risk below is grounded in an artefact in this folder or in a documented limitation; every control names the file or command that implements it.

Read with: docs/07-product-backlog.md (owners and Definition of Done), docs/08-test-plan.md (what is tested and what is not), docs/02-status.md and docs/04-modelling-decisions.md §7 (known limitations), workers/DEPLOYMENT.md (build, configuration and failure switches).

1. Risk register

Scoring. Likelihood and impact are H/M/L; exposure is the product on a 3-point scale (H=3, M=2, L=1), so exposure runs 1–9. A risk at exposure ≥ 6 needs a named contingency, not only a mitigation.

ID	Risk	Category	L	I	Exp	Mitigation (what is already in place)	Contingency / fallback	Owner (1st / 2nd)	Status	Trigger
R-01	The collaboration is too large for the Camunda 8 deployment batch ceiling: the whole BPMN resource is stored in every process definition record, so 15 processes mean 15 copies	Technical	H	H	6	v7.0 avoided this by keeping the four external suppliers as black-box participants; v14 has taken the risk on deliberately to connect them (docs/11-white-box-external-participants.md §5). Mitigations, in order: raise ZEEBE_MAX_MESSAGE_SIZE above the batch; deploy the nine hospital processes and the six outside participants in two batches; or set EXTERNAL_PROCESSES_EXECUTABLE = False in tools/spec_v2.py and rebuild so the six are drawn but not deployable. Verified against c8run 8.10.0-alpha5 : fifteen executable processes are rejected (nine records fit, the tenth is refused at 4,095,764 bytes). The shipped configuration implements the second mitigation — nine in the main file, the six outside participants in a 30 KB file — and both deployments succeed	Split the supplier work into a separate deployable file, or reduce diagram interchange payload; the walk-away is the smaller v1.0 collaboration with black-box suppliers	C / E	Mitigated and verified — see screenshots/static-checks-v14.txt §8. Residual: a further pool added to the main file without a second batch	Deployment rejected for size, or a new pool is added to the same file without a second batch

ID	Risk	Category	L	I	Exp	Mitigation (what is already in place)	Contingency / fallback	Owner (1st / 2nd)	Status	Trigger
R-02	A message throw has no correlation key, producing a live <code>JOB_NO_RETRIES</code> incident and a stalled pathway	Technical	M	H	6	The builder writes a <code>correlationKey</code> input mapping onto every throw event (fix for DEF-01); TC-WRK-06 verifies a blank key fails safely; both paths run with <code>incidents: none</code>	Re-deploy the last known-good resource; add the missing mapping by hand in the Modeler and re-run the path	A / C	Mitigated — fixed and re-tested	An incident named <code>JOB_NO_RETRIES</code> , or a message that never correlates
R-03	Correlation-key design lets a second hand-off for the same patient and purpose go uncorrelated, so a patient does not receive a second letter or funding decision	Technical / data	M	M	4	Purpose tags (<code>patientRef</code> + <code>"-referral-review"</code> , + <code>"-funding"</code> , + <code>"-refund"</code>) separate concurrent work; the key itself is the duplicate-suppression mechanism (docs/04-modelling-decisions.md §2)	Re-publish with a corrected key after cancelling the stale instance in Operate; if the purpose set is incomplete, add a tag and re-deploy	A / E	Monitored	A message is deliberately sent twice and only one instance is expected
R-04	Five external services are deterministic stubs, so passing tests are read as proof of real integration	Quality / technical	H	M	6	The simulation is stated in <code>workers/README.md</code> and repeated in <code>docs/08-test-plan.md</code> §5 and the limitations; forced-failure switches make the boundary branches genuinely testable	Replace one worker class per service when a real API exists; until then every claim is scoped to "simulated"	C / E	Accepted — documented limitation	A review or a marker treats a green run as real provider integration
R-05	Three modelled branches have never been driven end to end (declined payment, payment confirmation lost, three-month letter escalation)	Quality	H	M	6	Branches are modelled, the workers support them, and the switches to drive them are documented in <code>workers/DEPLOYMENT.md</code> §6	Drive each branch with <code>demo.payment.status</code> / <code>demo.pathway.scenario</code> and keep the log; if it cannot be done, state the gap at the review before a reviewer finds it	E / C	Open — carried as PB-30	A Sprint Review or the final demonstration asks for one of the three
R-06	Compensation is modelled and statically asserted but has never fired in a live run	Quality	M	M	4	Two compensation boundary events, two <code>isForCompensation</code> handlers, two compensation throws; the three worker types exist and register (docs/03-test-record.md §E3)	Force the giving-up path with <code>demo.external.unavailable-resources</code> and cap 3, and capture the handler completing in Operate	C / E	Open — carried as PB-31	A reviewer asks to see compensation work rather than read that it is wired
R-07	The Tasklist Task pane cannot draw the forms because <code>c8run</code> returns <code>401</code> from <code>/v2/authentication/me</code>	Technical / tooling	H	M	6	Forms are proven to render with <code>@bpmn-io/form-js</code> , the renderer Tasklist uses (<code>TOTAL=36 FAILED=0</code>), so the deliverable no longer depends on it (docs/03-test-record.md §F4–F5)	Demonstrate the form through the <code>form-js</code> render page and the Operate instance variables; treat Tasklist as an environmental gap, not a model gap	E / D	Open — carried as PB-32	A graded review requires a task completed in the Tasklist UI
R-08	No performance, load or concurrency testing has been done, and no target exists	Quality	M	M	4	Single-instance runs are deterministic and reproducible with <code>DemoClock</code> ; operation counts are small and the workers use 32 max jobs and 4 threads	If a load claim is needed, build a small parallel driver against the REST API; otherwise make no performance claim at all	E / C	Accepted — out of scope	A review expects evidence of behaviour under concurrent pathway load

ID	Risk	Category	L	I	Exp	Mitigation (what is already in place)	Contingency / fallback	Owner (1st / 2nd)	Status	Trigger
R-09	The accessibility expectation is not met: 0 grouping headings, only 71 descriptions across 230 form components, and no WCAG check	Quality	H	M	6	Field-level descriptions and validation are present; single-column layout; the gap was identified in <code>docs/01-gap-audit.md</code> §C3	Emit <code>group</code> components in <code>tools/forms_spec.py</code> for the long forms and re-run the form-js census; state the gap honestly until then	D / E	Open — PB-12 partly done	A marker inspects a long form or asks for the accessibility rationale
R-10	Two brief reduction targets remain unmet: 43 message start events against a target of ≤ 2 per pool, and 197 crossings against 42 in v1.0	Quality	H	M	6	The arithmetic that makes "under 80 events" unreachable without deleting scored elements is documented (<code>docs/05-v2.1-changes.md</code> §4); the crossing rise is confirmed by A/B measurement rather than asserted	Accept and explain, or spend the effort on the message-protocol change (PB-33) that would take start events to about 18; a crossing pass is PB-34	A / B	Accepted — documented trade	A reviewer measures the model and asks why the target was missed
R-11	Idempotency state lives in memory in one JVM, so a worker restart can allow a repeat that a persistent store would prevent	Technical / data	L	H	3	In-memory <code>HospitalStore</code> plus deterministic reference generation; the limitation is stated in <code>docs/04-modelling-decisions.md</code> §7 item 5	Restart the workers between runs and re-verify the reference; for a real deployment, move idempotency to database constraints on booking and payment references	C / A	Accepted — documented limitation	A worker JVM restarts mid-path, or the same patient is booked twice
R-12	Tooling versions are unusually specific and partly pre-release: c8run 8.10.0-alpha5, Camunda Modeler 5.51.0, <code>zeebe-client-java</code> 8.9.0 against an 8.10 engine, jar built under JDK 25 targeting Java 21	Technical	M	M	4	Versions are recorded in the test record and the plan; the jar targets Java 21 so it runs on a 21 JVM; the model declares <code>Camunda Cloud / 8.10.0</code> and exporter 5.51.0	Pin every version in writing; if the engine moves, re-run the two paths and the preservation audit before claiming anything; a mismatch in an alpha feature is the first thing to suspect	C / E	Monitored	An engine or Modeler upgrade, or a client/engine incompatibility at start-up
R-13	Documentation drifts from the artefact: one document described a diagram state that was never shipped; the worker README's job-type table is three types short; one cited screenshot does not exist	Quality / CM	H	M	6	The preservation audit and the parser gate compare the file with its predecessor rather than with prose; the drift items are logged as DEF-14 to DEF-17	Correct the docs or complete the pass before submission. Executed in R3: the drifted diagram document was deleted and replaced by <code>docs/06-diagram-engineering.md</code> , and the <code>workers/README.md</code> job-type catalogue was completed and its counts corrected. DEF-14, DEF-15 and DEF-16 are closed; DEF-17 (one cited screenshot that does not exist) stays open	E / A	Mitigated — three of four drift items closed, the fourth tracked	A reviewer opens the worker README or a cited screenshot beside the model

ID	Risk	Category	L	I	Exp	Mitigation (what is already in place)	Contingency / fallback	Owner (1st / 2nd)	Status	Trigger
R-14	There is no version-control repository in the workspace, so a change can silently overwrite another member's work and there is no branch history	CM / schedule	M	H	6	The shared workspace is the integration point; the numbering convention makes the latest document obvious; the generator is deterministic and the preservation audit catches a lost scored element; hand-overs are logged in <code>docs/07-product-backlog.md</code> §7	Initialise a repository at the next session and commit the current state as the baseline; until then take a timestamped folder copy before any bulk regeneration	E / A	Open	Two members edit the model or a document in the same session
R-15	Single point of failure: one member holds the layout engine, one holds the worker catalogue, and one holds the generator	People	M	H	6	Every backlog item has a second owner; the hand-over log records ten explicit hand-overs with an acceptance step; the tools are committed in <code>tools/</code> and documented in <code>README.md</code>	The second owner re-runs the relevant gate and takes the item; the generator is deterministic so the output can be reproduced without the original author	A / B / C (each is 2nd owner for the others)	Mitigated by the two-owner rule	The first owner is unavailable at a graded stand-up or review
R-16	Cross-team hand-offs are ambiguous: the case study does not say who owns some tasks, and two teams share a form	People / quality	M	M	4	Candidate groups match the pool roles (39/39); shared forms are deliberate and checked by the build; the ambiguity in the case study is recorded as six numbered assumptions in <code>docs/04-modelling-decisions.md</code> §6	Where the case study is silent, present the assumption as an assumption; do not defend it as a requirement	A / D	Accepted — documented	A reviewer disputes which team owns a task
R-17	The system processes sensitive personal, clinical and financial data	Data protection	L	H	3	Test data is synthetic (<code>PAT-DEMO-...</code>); the payment workers never read or emit card data; audit workers do not write clinical variables back; no real provider is contacted	If any real data is ever introduced, stop and handle it under the hospital's own policy; for this portfolio, keep the synthetic-data rule and say so	C / E	Accepted with the synthetic-data control	Any real patient identifier appears in a form, log or screenshot
R-18	A rollback is misunderstood: Camunda 8 keeps every deployed version and cannot delete one, so "rolling back" means re-deploying the earlier file, which creates a further version	Technical / CM	M	M	4	Deployment versions are recorded (v17 processes, v18 forms after the v2.1 rebuild); the previous files exist in the sibling folders; instances are cancelled rather than deleted	Re-deploy the previous resource, then cancel any instances created by the faulty version in Operate; tell the audience the version number being shown	E / C	Understood — mechanism documented here	A live demonstration shows a defect and the previous resource is needed
R-19	The timebox is missed and the release is not finished	Schedule	M	H	6	MoSCoW with explicit carry-over; the Must list is the release; PB-30 to PB-35 were held as stretch items and not started, so no committed item was dropped	Submit the last complete held release. <code>docs/02-status.md</code> records that v1.0 is untouched, still deploys, still runs both paths, and has its own evidence, so v1.0 is the safer submission if v2.x cannot be finished	E / A	Monitored	A committed item is still In progress at the review

ID	Risk	Category	L	I	Exp	Mitigation (what is already in place)	Contingency / fallback	Owner (1st / 2nd)	Status	Trigger
R-20	The engine or cluster is unavailable at the demonstration	Technical / schedule	M	H	6	Both scripted paths are deterministic and re-runnable; logs and Operate screenshots are retained in <code>screenshots/</code> ; <code>DemoClock</code> makes runs reproducible	Restart <code>c8run</code> , re-run <code>tools/deploy.sh</code> , restart the workers, then re-run the path; if the engine will not start, present the retained log and screenshots and say plainly that the run is recorded rather than live	E / C	Mitigated by retained evidence	<code>GET /v2/topology</code> does not return 200, or the deployment is rejected
R-21	No data-protection impact assessment exists, although the case study expects the information and authorisation requirements to be analysed	Data protection / quality	M	M	4	The data-protection rules that are implemented are stated and testable: no card data (TC-NFR-01), synthetic test data (TC-NFR-08), role separation (TC-NFR-02), non-overwritable audit records (TC-NFR-03)	Produce a short DPIA-style note over the implemented controls; if it is not produced, present the controls as controls and do not claim compliance	E / A	Open	A reviewer asks how the data-protection requirement was analysed

Exposure summary. Twelve risks sit at exposure 6 and each therefore has a named contingency, not only a mitigation: R-01, R-02, R-04, R-05, R-07, R-09, R-10, R-13, R-14, R-15, R-19 and R-20. Of those, R-04, R-05, R-07, R-09, R-10 and R-13 are open quality or tooling risks carried into submission, and R-14 (no version-control repository) is an open configuration-management risk. Everything else is either mitigated in the artefact or accepted as a documented limitation.

2. Contingency plan for the demonstration

The demonstration is the risky moment: a live engine, a Java process, a large model and a browser all have to cooperate in one timebox. The plan names what to do when a component fails, in the order a presenter would hit it.

Failure at the demonstration	Immediate action	Fallback evidence to show instead	Owner
The laptop or <code>c8run</code> will not start	Start the bundled <code>c8run</code> ; if it fails within 5 minutes, stop trying	Retained <code>screenshots/run-happy-path.log</code> and <code>run-exception-path.log</code> , plus the Operate screenshots	E
<code>GET /v2/topology</code> does not return 200	Restart <code>c8run</code> and re-check once	As above; do not spend the slot debugging infrastructure	E
Deployment rejected	Re-run <code>tools/deploy.sh</code> once and read the error; check resource size (R-01)	Show the last accepted deployment evidence (<code>docs/03-test-record.md</code> §C1, §H2a)	C
Workers fail to start or subscribe	Re-run <code>java -jar workers/target/hospital-external-workers-1.0.0.jar</code> ; check the gateway address	Show the retained subscription log and explain the 37 subscriptions	C
The path stalls at a user task	Do not skip the task manually; check the incident search first	Switch to the exception path, which is shorter, or present the retained log	E

Failure at the demonstration	Immediate action	Fallback evidence to show instead	Owner
An incident appears live	Capture it, name it honestly, then cancel the instance in Operate	Use the incident as the demonstration of the exception design rather than pretending it did not happen	E
Tasklist will not draw a form	Do not debug authentication live (R-07)	Show <code>screenshots/camunda-forms-rendered.png</code> and the form-js census	D
The diagram is too large to navigate	Use <code>diagram/sections/</code> and the per-pool crops	Show one pool crop at a time	B
Time runs out	Stop adding material; close with the limitations	The limitations list in <code>docs/04-modelling-decisions.md</code> §7 is the closing script	A

Presenting rule. A recorded result is always labelled as recorded. A simulated service is always called simulated. A branch that was not driven is never implied to have run. This is the same rule the test plan applies in `docs/08-test-plan.md` §7.

3. Rollback plan

Camunda 8 keeps every deployed version and does not delete one, so rollback is a re-deploy of the previous resource plus cancellation of the instances the faulty version created. The mechanism is stated here so that it cannot be discovered for the first time during a demonstration.

Component	Roll forward normally	Roll back to the previous release
Model and forms	<code>python3 tools/build_v2.py</code> then <code>bash tools/deploy.sh</code>	Copy the previous <code>.bpmn</code> and <code>model/forms/</code> from the sibling folder (<code>../UFCEP6-0-3_BPMN_Hospital_Referral_v2/</code> for v2.0, <code>../UFCEP6-0-3_BPMN_Hospital_Referral/</code> for v1.0), then deploy; the engine records a further version rather than replacing one
Running instances	Let the path complete	Cancel the affected instances in Operate; they are not deleted and remain visible as cancelled
Workers	<code>mvn -q clean package</code> and restart	Stop the JVM and restart the previous jar; the in-memory store resets, which is also the reset mechanism for idempotency tests
Demo state	Optionally pin <code>demo.clock.today</code> and the failure switches in <code>application.yaml</code>	Set the switches back to their defaults (<code>simulated-failure.enabled: false</code> , <code>payment.status: APPROVED</code> , <code>force-no-slots: false</code> , <code>unavailable-resources: []</code>)
Diagram and sections	<code>bash tools/render_diagram.sh</code> and <code>python3 tools/export_sections.py</code>	Re-run both against the restored model; they are generated artefacts and hold no state
Release as a whole	Deploy the current tag	Submit the previous held release; <code>docs/02-status.md</code> records that v1.0 still deploys and runs both paths, and that the v1.0 folder is untouched

Decision rule. Roll back if a Must item fails at Sprint Review 2 and cannot be fixed within the review slot. State the rollback and the reason; do not present an unverified release as if the rollback had not happened.

4. Configuration-management plan

4.1 Repository layout

The shared workspace is this folder. There is no `.git` directory in it (checked), so the layout and the conventions below, plus the deterministic build, are what keep the configuration under control. See Assumption A1.

Path	Contents	Under configuration control?
<code>model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn</code>	The single deployable collaboration	Yes — generated, committed
<code>model/forms/*.form</code>	36 Camunda Forms	Yes — generated, committed
<code>workers/</code>	Java sources, <code>pom.xml</code> , <code>run-workers.sh</code> , <code>README.md</code> , <code>DEPLOYMENT.md</code> , <code>src/main/resources/application.yaml</code>	Yes — hand-written
<code>tools/</code>	Generator (<code>build_v2.py</code> , <code>spec_v2.py</code> , <code>bpmn_builder_v2.py</code> , <code>layout_engine.py</code> , <code>forms_spec.py</code>), <code>deploy.sh</code> , <code>demo_scenario.py</code> , <code>render_diagram.sh</code> , <code>export_sections.py</code> , <code>analyse_layout.py</code> , <code>verify_preservation.py</code> , <code>validate_model.js</code>	Yes — hand-written
<code>docs/NN-topic.md</code>	The working notes, this plan and the reviews	Yes — hand-written
<code>README.md</code> , <code>.bpmnlintrc</code>	Entry point and lint configuration	Yes
<code>screenshots/</code>	Run logs and UI evidence	Yes for the logs; the PNGs are capture output
<code>diagram/</code> , <code>diagram/sections/</code>	SVG, PNG and per-pool crops	Generated — regenerable, treated as release evidence

4.2 Revision and release scheme

Release	Model file	Design document	zeebe:versionTag in the file	Engine deployment version observed
v1.0	<code>../UFCEP6-0-3_BPMN_Hospital_Referral/model/...v2.bpmn</code> (untouched)	—	—	—
v2.0 (Initial Release, R1)	<code>model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn</code>	<code>docs/01</code> – <code>docs/04</code>	<code>2.0.0</code> on all 9 processes	—

Release	Model file	Design document	zeebe:versionTag in the file	Engine deployment version observed
v2.1 (second release, R2)	same path, regenerated	docs/05-v2.1-changes.md	2.0.0 — not bumped	processes at v17, forms at v18

The folder is the unit of release; each release keeps its own copy of this workspace, and the previous releases remain deployable. The model itself carries `versionTag` 2.0.0 on every process, which is a genuine shortcoming for R2: a drawing release that changes the committed artefact should have carried a distinct tag. It is recorded as an open configuration-management action (feedback FB-10 in `docs/10-sprint-reviews-and-feedback.md`), and the engine's own deployment versions (v17 / v18) are the working identifier in the meantime.

4.3 Naming and identifier conventions

Thing	Convention	Examples
Participants	P_<Team>	P_MedicalSecretaries , P_TreatmentBookings , P_ExternalScheduling
Executable processes	kebab-case pool name	medical-secretaries , treatment-bookings
Lanes	<pool>_L<n>	consultants_L0 , finance-team_L1
Elements by kind	<POOL>_<kind>_<Name>	SEC_Task_CheckPack , TRT_Auto_ProcessPayment , CON_Bnd_LetterOverdue , OUT_EGW_... , SEC_EGW_InfoRequest , TRT_Comp_ReleaseSeries , SUB_Pathway_Report
Camunda Forms	kebab-case, action-first	check-referral-pack , determine-funding-route , resolve-dispatch-problem
Job types	<domain>.<verb-noun>	referral.check-supporting-documents , payment.process-transaction , treatment.release-series
BPMN errors	Error_<CODE> with errorCode="<CODE>"	Error_REFERRAL_PACK_UNREADABLE / REFERRAL_PACK_UNREADABLE
Messages	Message_<event>	Message_referral_received , Message_letter_escalation_higher_management
Documents	NN-topic.md , two-digit prefix	07-product-backlog.md
Backlog / test / defect / risk / hand-over IDs	PB-nn , TC-<AREA>-nn , DEF-nn , R-nn , H0-nn	PB-26 , TC-BPMN-12 , DEF-15 , R-13 , H0-09

The conventions are enforced where they can be: `tools/validate_model.js` fails a user task without a form or candidate group and a service task without a job type; `tools/build_v2.py` fails a form that is bound but not defined, defined but not bound, or two elements sharing a grid cell; `tools/verify_preservation.py` compares element ids, job types, error codes and candidate groups against the previous release.

4.4 What is generated and what is authored

Generated artefact	Authored source	Command
model/...v2.bpmn	tools/spec_v2.py , tools/bpmn_builder_v2.py , tools/layout_engine.py	python3 tools/build_v2.py
model/forms/*.form	tools/forms_spec.py	python3 tools/build_v2.py
diagram/hospital-patient-pathway-v2.svg and .png	the model	bash tools/render_diagram.sh
diagram/sections/*.png	the model	python3 tools/export_sections.py

Everything else — workers, tools that are not the generator, documents — is authored. The authoring rule is that a generated file is never hand-edited: a change is made in the source and the artefact is regenerated, which is what makes the v2.0-to-v2.1 diff reviewable. docs/05-v2.1-changes.md §5 states that build_v2.py is deterministic and produces byte-identical output for the same inputs, so an unexplained diff is a real change to be reviewed.

4.5 Keeping the deployable model and its forms in step

Camunda 8 stores the whole BPMN resource inside each process definition record, so the model and its forms must go up in the same deployment. Three controls hold this together:

1. **One batch, one command.** tools/deploy.sh collects every model/forms/*.form and the single BPMN file into one multipart deployment, so a form can never lag its model on the cluster by being uploaded separately.
2. **Build-time binding check.** tools/build_v2.py cross-checks every form bound by a zeebe:formDefinition against the form definitions and refuses to emit when either side is missing; the parser gate independently confirms 39/39 user tasks have a form and a candidate group.
3. **Regeneration regenerates both.** The same command writes the BPMN and all 36 forms from the same sources, so the model and the forms cannot be regenerated from different revisions without the diff showing it.

4.6 Change review and integration

A change is integrated only when all of the following hold; this is the same gate set as the Definition of Done in docs/07-product-backlog.md §6.

Gate	Command	Failure action
Generator accepts the change	python3 tools/build_v2.py	Fix the source; never edit the generated file
Layout rules hold	python3 tools/analyse_layout.py model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn	Re-run the layout pass; overlaps must be 0
Nothing scored was lost	python3 tools/verify_preservation.py	Restore the element before anything else is accepted
The Modeler's parser is clean	node tools/validate_model.js model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn	Fix the model; 0 warnings and 0 unresolved references

Gate	Command	Failure action
The release still deploys	<code>bash tools/deploy.sh</code>	Fix or roll back (§3)
The behaviour still runs	<code>python3 tools/demo_scenario.py</code> and <code>--exception</code>	Re-run after the fix; a pass recorded against an undeployed model is discarded
A second owner has read it	hand-over log, <code>docs/07-product-backlog.md</code> §7	The item is not Done

4.7 Reproducing the whole release from source

```
# 1. regenerate the model and the 36 forms
python3 tools/build_v2.py

# 2. prove nothing scored was lost by the regeneration
python3 tools/verify_preservation.py

# 3. measure the drawing
python3 tools/analyse_layout.py model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn

# 4. parse it the way the Modeler does
NODE_MODULES_DIR="$PWD/tools/node_modules" node tools/validate_model.js \
    model/UFCEP6-0-3_Hospital_Patient_Pathway_v2.bpmn

# 5. build, deploy and run
cd workers && mvn -q -o clean package && cd ..
bash tools/deploy.sh
java -jar workers/target/hospital-external-workers-1.0.0.jar &
python3 tools/demo_scenario.py
python3 tools/demo_scenario.py --exception

# 6. regenerate the release evidence
bash tools/render_diagram.sh
python3 tools/export_sections.py
```

5. Assumptions

- A1 — No version control.** There is no `.git` directory or other VCS metadata in this folder or its parents (checked). This plan therefore describes configuration management as it exists: a shared folder, numbered documents, a deterministic generator, three executable gates and a two-owner review. Adopting a repository is recorded as feedback FB-12 and risk R-14.

2. **A2 — Branching.** No branching or tagging history exists to describe. The scheme in §4.2 is the *effective* model (each release is a held copy of the folder, integrated only through the gates); it is not a record of a branching strategy in use.
3. **A3 — Version tag not bumped.** The `zeebe:versionTag` in the shipped file is `2.0.0` on all nine processes even though the folder was regenerated as v2.1. This is stated as a finding rather than presented as intended versioning; the engine deployment versions (v17 processes, v18 forms) are the identifiers actually observed.
4. **A4 — Release identity.** "Initial Release" is R1 (v2.0) and the "second release" is R2 (v2.1), per `docs/07-product-backlog.md` §2.
5. **A5 — Risk owners.** Owners are planning assignments to team members A–E using the roles in `docs/07-product-backlog.md` §4; no individual is named in any artefact.
6. **A6 — Exposure scores.** Likelihood and impact are the authors' judgement against the evidence in this folder. They are not derived from incident data, because no incident history is retained beyond what `docs/03-test-record.md` records.
7. **A7 — Documentation drift.** One cited screenshot does not match the shipped artefacts (defect DEF-17 in `docs/08-test-plan.md` §9). DEF-14, DEF-15 and DEF-16 were closed in R3 by completing the `workers/README.md` catalogue and replacing the drifted diagram document with `docs/06-diagram-engineering.md`. This plan records the remainder as risk R-13 and feedback items in `docs/10-sprint-reviews-and-feedback.md` §4.